

Operating System Principles

Introduction and Overview

Assumed background

- Algorithms and data structures
- Comfortable programming in C/C++
- Preliminary understanding of computer architecture
 - We will cover some basics in this course

Before we begin ...

- Some advice
 - Speak up in class, ask questions
 - Attend all classes
 - Hint: Ideas concerning projects and exams
 - Do all projects on your own!
 - Do not copy or lift code from other sources
 - Take notes in class
 - Read text-book soon after class
 - Even better: read before class and come prepared with questions

Operating Systems: Introduction & Background

What is an operating system?

- Let us begin with an incomplete definition

Software that allows multiple programs to run on the same computer

- Why is such software needed?

Example Scenario

- Imagine a laptop running a browser and an music player

- Browser code (fake):**

...
load R1, @100
load R2, @200
add R3, R1, R2
...

- Player code (fake):**

...
load R1, @700
load R2, @900
sub R3, R1, R2
...

- Can you identify some of the resources that both these programs are using (i.e., sharing)?
 - CPU (including registers R1-R3), main memory, busses
- Could this cause a problem?

Example Scenario

- Consider the following “interleaving” of instructions

...

load R1, @100

load R1, @700

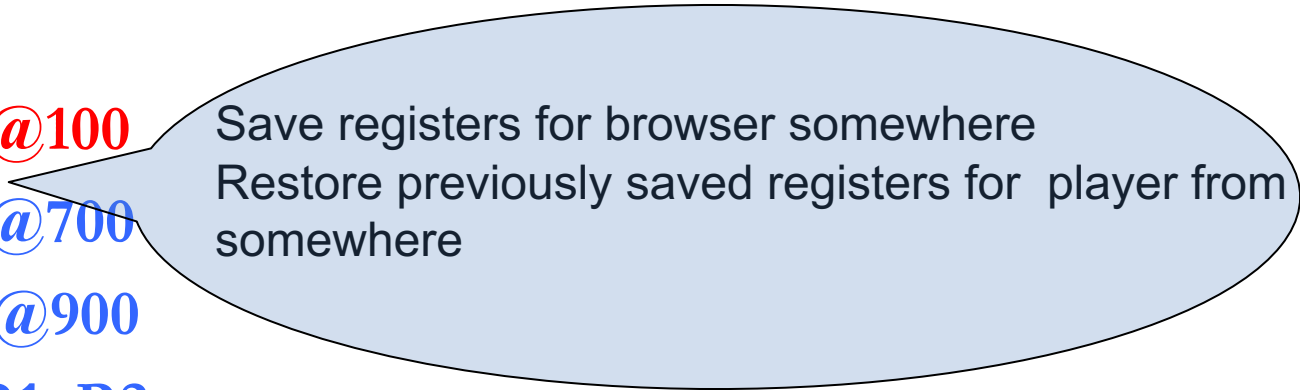
load R2, @900

sub R3, R1, R2

load R2, @200

add R3, R1, R2

...



Save registers for browser somewhere
Restore previously saved registers for player from somewhere

- What's the problem here?
- What should happen?

Questions, questions, ...

- Sharing of CPU
 - Where would the contents of the registers be saved?
 - Answer: Main memory (where data and instructions are also stored)
 - By whom?
 - When should the CPU decide to stop executing the browser and start executing the player?
- What about similar problems due to sharing of other resources – busses, h/w caches, memory, IO devices?

Option #1: Program-driven Sharing

- Let the programs take charge of dealing with sharing-related problems
 - E.g., for CPU registers: Each program saves its registers in some memory locations before the next starts executing; when it is resumed, it loads registers from these memory locations
- Will this work?
- Answer: No, unless care is taken
 - E.g., a program overwrites the memory holding registers for another

Option #1: Program-driven Sharing

- Let the programs take charge of dealing with sharing-related problems
 - E.g., for CPU cycles: Each program executes instructions after which it gives up control to another program
- Will this work?
- Answer: No!
 - What if a program is malicious and doesn't yield the CPU -> others will “starve”

Option #2: Hardware-driven Sharing

- Let the hardware take charge of dealing with sharing-related problems
 - E.g., for CPU cycles and registers: Let the CPU itself decide when to run which program and save/restore registers
- Will this work?
- Answer: Yes, but ...
 - We are stuck with the policy (e.g., scheduling algorithm) implemented in hardware
 - Desirable for some resources (e.g., busses) but not for all

Summary

- Sharing of resources among can cause problems related to
 - Correctness
 - E.g., a program's register contents being corrupted by another
 - Fairness
 - E.g., a program not getting enough CPU cycles
- Program-driven solutions don't suffice; Hardware-driven solutions likely to be inflexible or complex/costly
- Need for a software other than the programs that facilitates correct and fair sharing of hardware resources
 - **The operating system**



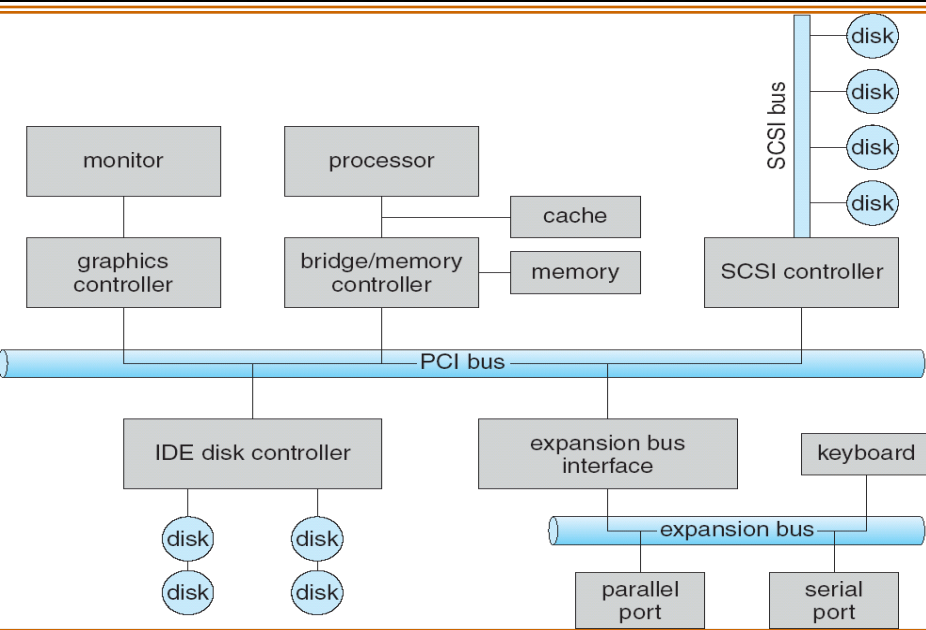
Programs

- Sharing of resources could be done by h/w, OS, or even programs

- E.g., Sharing of busses done by h/w
- E.g., sharing of CPU cycles done by OS
- E.g., sharing of registers done by programs themselves



OS



Hardware

- Discuss other resources and their sharing

