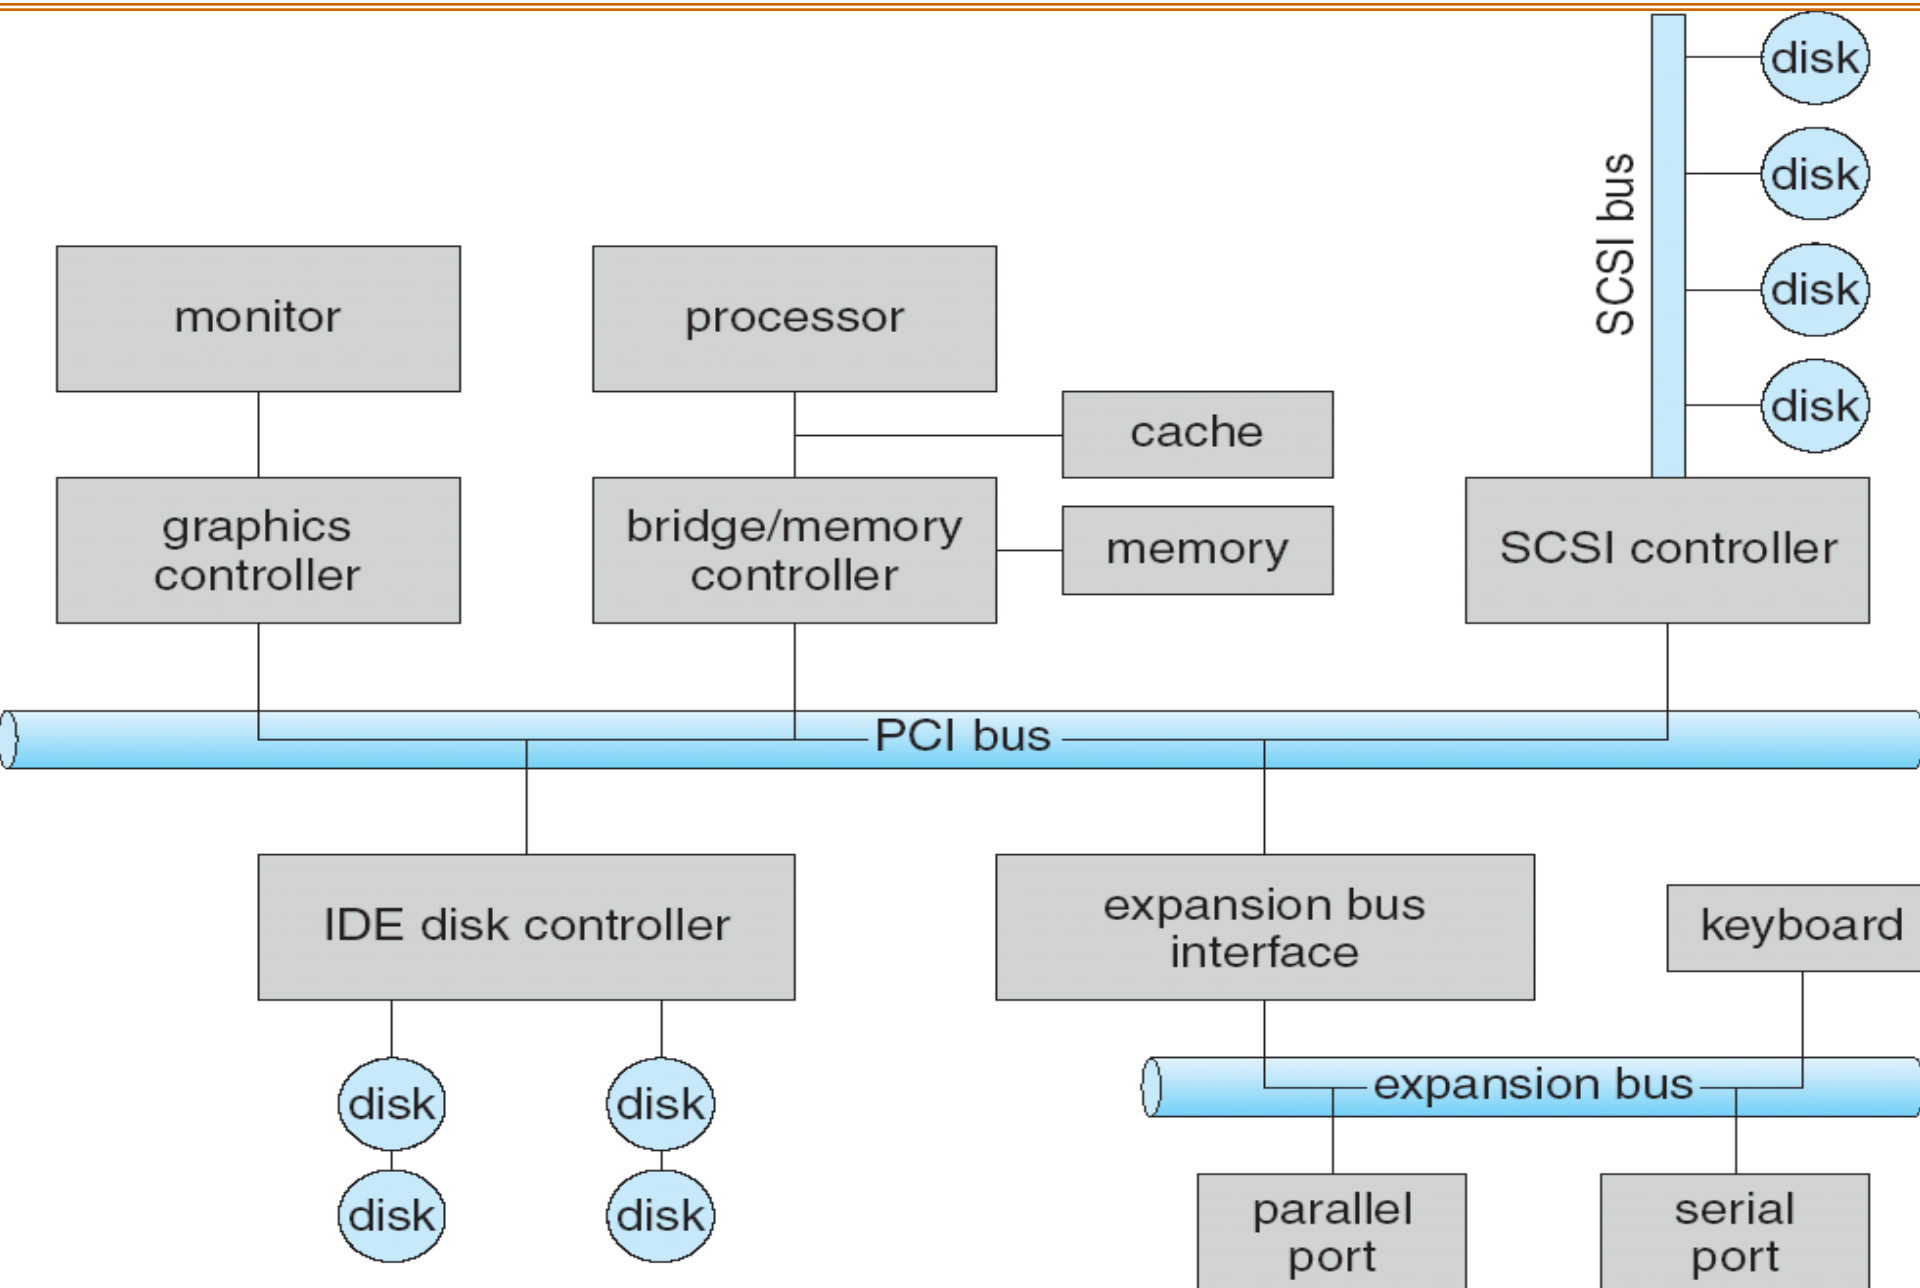


Operating Systems Principles

Recap of Some Basics



Summary

- Sharing of resources can cause problems related to
 - Correctness
 - E.g., a program's register contents being corrupted by another
 - Fairness
 - E.g., a program not getting enough CPU cycles
- For certain resources OK to have hardware-driven sharing
- For others, would like to have software-driven sharing

Summary

- Need for a software **other than the programs** that facilitates correct and fair sharing of hardware resources
 - **The operating system**
 - May require special support from H/W (e.g., memory protection for register sharing that we saw in last class)



Programs

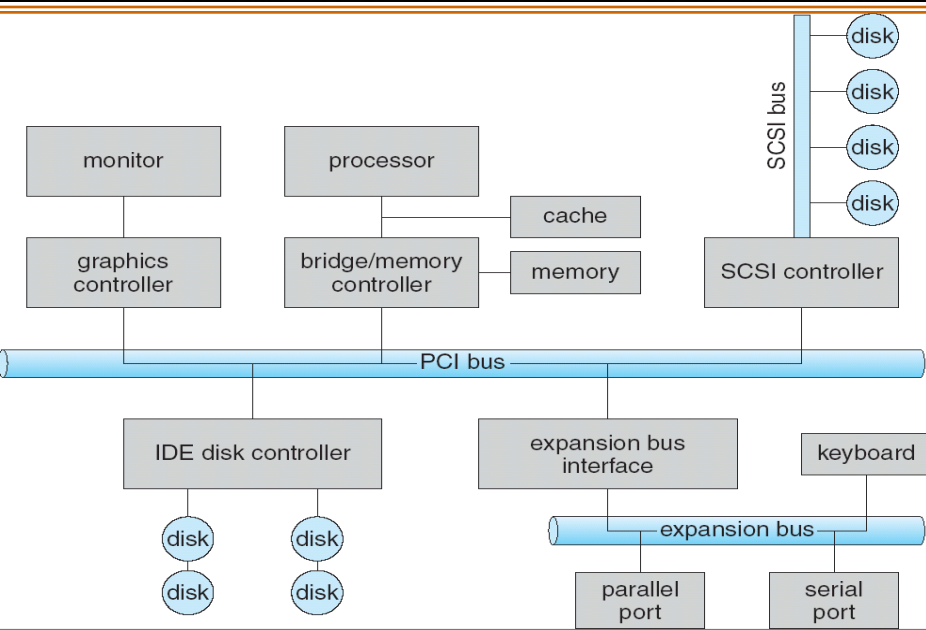
- Sharing of resources could be done by h/w, OS, or even programs

- E.g., Sharing of busses done by h/w
- E.g., sharing of CPU cycles done by OS
- E.g., sharing of registers done by programs themselves

OS

- Discuss other resources

Hardware and their sharing



A More Complete Definition

Software that offers following functionality

- OS is a **resource allocator**
 - Divides up resources among programs
 - Decides between conflicting requests for efficient and fair resource use
- OS is a **control program**
 - Controls execution of programs to prevent errors and improper use of the computer
- OS **virtualizes** resources
 - Hides hardware details and offers programs easier interfaces to work with

What We Will Study

- How the OS allocates, controls, and virtualizes resources to ensure safety (isolation), performance, and fairness

- Before diving into OS design, we need to understand some basics
- Certain aspects of CPU operation
 - Instruction Set Architecture (ISA)
 - Instruction execution
 - CPU modes
 - Interrupts and Traps
- Certain basics about programs
 - Compilation, Linking, and Loading
- Certain aspects of systems programming
 - Using system calls

Recap of Basics

- Key to OS acting as a control program and a resource controller are following CPU/memory facilities
 - #1: Multiple CPU modes and privileged instructions
 - #2: The notion of traps (part of CPU design)
 - Related concepts: Interrupts and signals
 - #3: Memory address translation and protection facilities (hardware called memory controller)
 - Related: Translation Lookaside Buffer (TLB)

ISA

- The set of instructions the CPU offers
 - Interface provided by CPU to software
 - That is, to both programs and OS
- Examples
 - Load (from memory to register)
 - Store (from register to memory)
 - Jump
 - ALU instructions (add, sub)
 - Halt (the computer)
 - IO instructions, and many more

Instructions Involving Registers and Memory

- Generic examples
 - Load Reg, @MemAdd
 - Store Reg, @MemAdd
- Add Reg1, Reg2, Reg3

The Need for “Privileged” Instructions

- Can a program execute all instructions in the ISA?
 - Should it?
- Answer: No
- Why?
- E.g., what could go wrong if a program were allowed to execute the “halt” instruction?
- E.g., what could go wrong if a program was allowed to execute IO instructions? (will become more clear later in the course)

Privileged Instructions

- **Key Idea #1:** Let some (appropriate) subset of ISA be privileged
 - Only the OS gets to execute privileged instructions
- The CPU is designed in a way to disallow programs to execute privileged instructions
- But how would the CPU know the difference between a program and the OS?
 - An instruction is an instruction!

Dual CPU Mode

- **Dual-mode** operation allows OS to protect itself from programs and programs from each other
 - **User mode** and **OS/kernel mode**
 - **Mode bit** provided by hardware
 - Provides ability to distinguish when system is running user code or kernel code
- If user code executes priv. instruction, CPU enters a special error-like state and control jumps to OS

Dual CPU Mode

- OS runs with CPU in kernel mode
- Is responsible to ensure programs run with CPU in user mode
- What is required to realize the above?
 - OS is the first software to run!
 - The booting up of the OS
 - OS has the ability to change CPU mode from kernel to user
 - Programs have the ability to change CPU mode from user to kernel

Traps

- **Key Idea #2:** Let the CPU be designed so it switches to kernel mode whenever certain conditions that require OS attention occur
 - These special conditions are called traps
 - Example 1: Programs are offered a special instruction via which they can raise a trap
 - Used to implement system calls
 - Example 2: Segmentation fault, division by zero, more to come

Traps

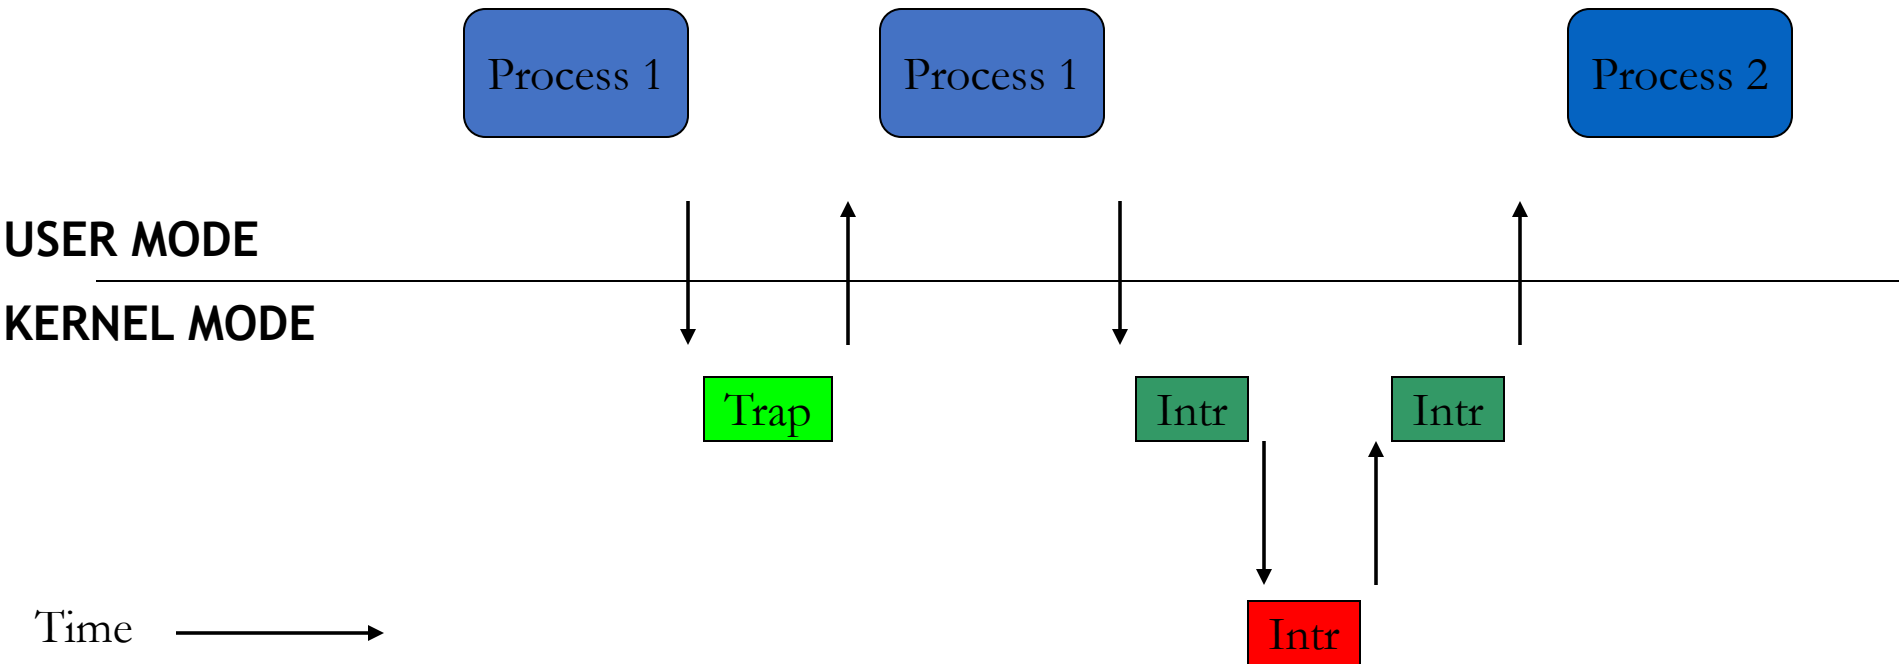
- On detecting trap, CPU must:
 - Save process state
 - Transfer control to trap handler (in OS)
 - CPU indexes *trap vector* by trap number
 - Jumps to address
 - Restore process state and resume

0:	0x00080000	Illegal Address
1:	0x00100000	Memory Violation
2:	0x00100480	Illegal Instruction
3:	0x00123010	System Call
...	...	

Interrupts

- Special conditions **external to the CPU** that require OS attention
 - Note difference from traps
- CPU designed to switch to kernel mode upon detecting an interrupt
- Example: A keystroke raises an interrupt

Interrupts and Traps



- Only two ways to enter kernel mode from user mode