


Operating Systems Principles

CPU Management (Processes)

How Does the OS Help Multiple Processes Share the Same CPU?

E.g., Two processes on a CPU

- Lets consider (only) two processes A and B that are running on the same CPU
- Let us look closely at some illuminating events in such a system




How does the OS get to start running here?

- Recall: user- \rightarrow privileged mode occurs only when a trap is raised or an interrupt occurs
- Can't rely on traps – why?


The Role of Interrupts

- There must be an interrupt mechanism via which the OS gets a chance to run on the CPU every so often
 - A timer interrupt that periodically lets the OS run, typically, once every few milliseconds



How does the OS decide which process to run next?

- Done by the CPU scheduler within the OS
- Will study later



How do we ensure that A resumes execution at T2 as if it had not been taken off the CPU at T1?

- By ensuring that we save the entire “state” of A at T1 and can resume it from this state at T2
- $\text{state}(A, T1) == \text{state}(A, T2)$
- What does the state of A at T1 consist of?

State of A at time $T1$ (1)


- #1: Contents of A 's address space
 - What the code, data, heap, and stack contain at $T1$

Contents of A's Address Space


0xFFFFFFFF

virtual addresses

0x00000000




stack



heap



data



code

State of A at time T1 (1)

- #1: Contents of A's address space
 - What the code, data, heap, and stack contained at T1
- Q: Where do these contents reside at time T1?
 - In portion of main memory set aside for A
 - We rely on memory manager to ensure they remain unchanged by other processes during [T1, T2]
 - More details when we study virtual memory management

State of A at time $T1$ (2)


- #2: Layout of A 's address space
 - The address ranges the code, data, heap, stack span

Layout of Address Space

0xFFFFFFFF

virtual addresses

0x00000000

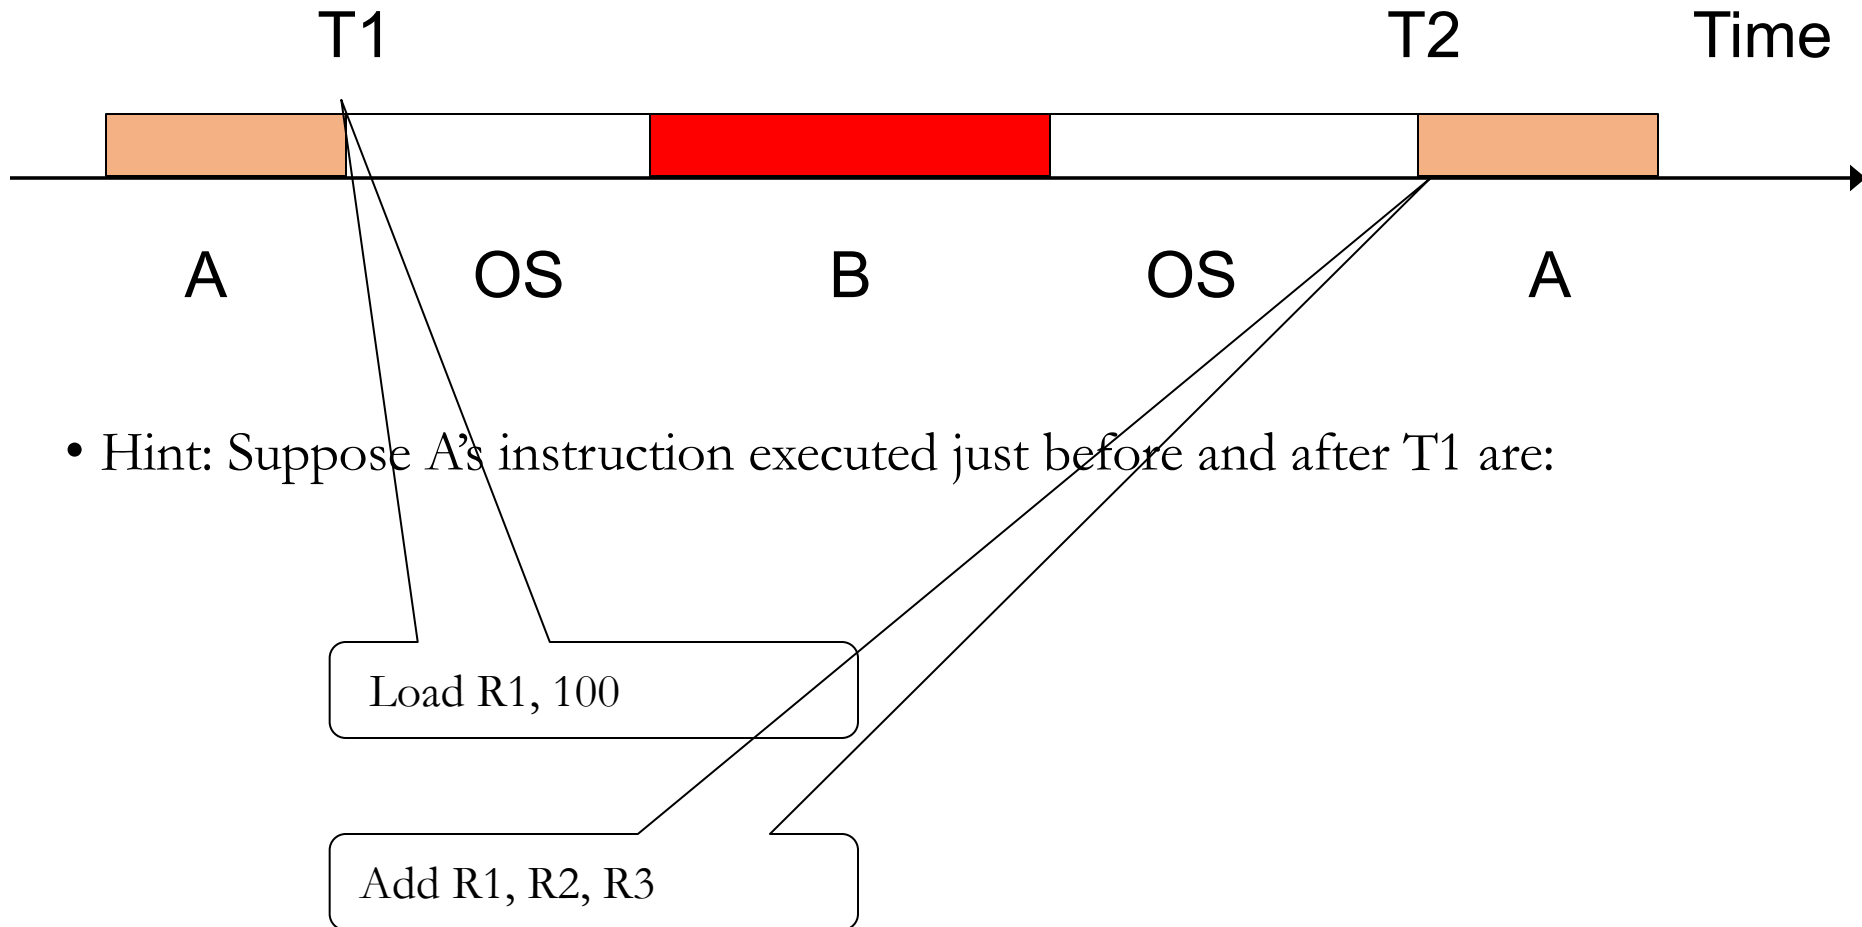


State of A at time T1 (2)

- Layout of A's address space
 - The address ranges the code, data, heap, stack span
- Q: Where are these address ranges stored?
 - Somewhere in memory
 - In whose address space? Lets say A's own address space for now (e.g., its heap or data segment)
 - **OS will need to maintain at least one address, say where A's address space begins in main memory**

State of A at time T1 (3)

- Anything else?




State of A at time T1 (3)

- #3: All the register values at time T1 need to be saved in main memory and restored at time T2
- Called the **hardware context** of process A
- Typically, the hardware context also has information about address space layout
 - E.g., Stack Pointer Register
 - Why does the CPU provide this facility?
 - Because registers are accessed much faster than memory

Process Control Block (PCB)

- A data structure in the operating system's address space where it maintains a variety of information about the process
 - Example we learnt today: Values held in various registers when the process was context switched out

Context Switch



```

context_switch() {

    push R0, R1, ...           // save regs on its stack
    PCB[curr].SP = SP          // save stack pointer
    PCB[curr].PT = PT          // save ptr(s) to address space

    next = schedule()          // find next process to run

    PT = PCB[next].PT
    SP = PCB[next].SP
    pop Rn, ... R0

    return                      // NOTE: Ctrl returns to another process
}

```

Overheads of Process Switch

- Direct
 - “Wasted” instructions of outgoing process
 - The time spent on switching context
- Indirect
 - Cache pollution
 - TLB flush