


Operating Systems Principles

CPU Management (Processes)

How Does the OS Help Multiple Processes Share the Same CPU?

Overheads of Process Switch

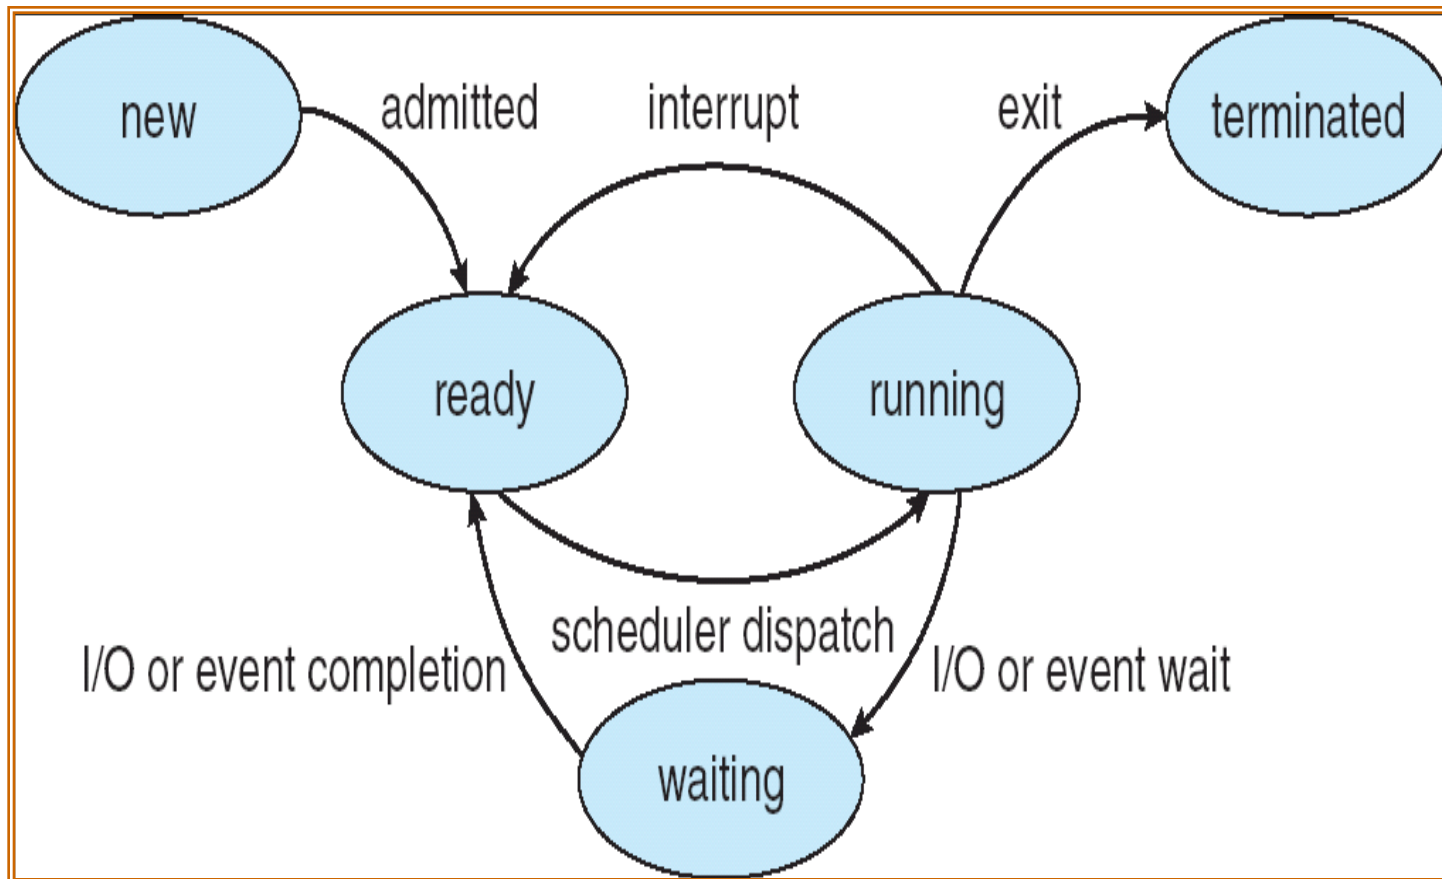
- Direct
 - “Wasted” instructions of outgoing process
 - The time spent switching context
- Indirect
 - Cache pollution
 - TLB flush



How does the OS decide
which process to run next?

- Done by the CPU scheduler within the OS
- Today's topic

Not All Processes May be Willing/Able to “Run”



Process Control Block (PCB)



- A data structure in the operating system's address space where it maintains a variety of information about the process
 - Example of information we learnt in last class: Values held in various registers when the process was context switched out
 - Another important piece of information held by the OS in a PCB: what the scheduling-related “state” of a process is (new, ready, running, waiting, terminated)

Some data structures that group subsets of PCBs together


- Queue of ready (runnable) processes
 - Scheduler picks from these
- Queues of waiting (blocked) processes
 - Separate queues for different devices or software resources (e.g., locks)

Some food-for-thought related to PCBs

- Sometimes PCB of exited process kept in memory
 - Pop quiz: Why?
- All PCBs in a part of RAM reserved by the kernel for itself and inaccessible to processes
 - Why?
 - This part of RAM initialized during boot-up





Running



Ready


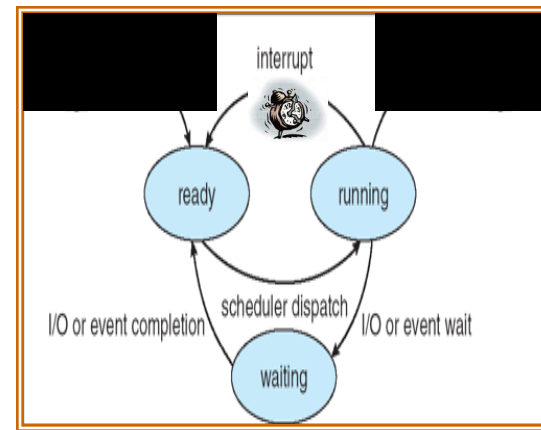
Lock

Disk






Waiting

Timer interrupt




Lock






Running



Ready




Lock

Disk




Waiting

I/O system call





Running




Ready


Lock



Waiting



Lets pick the second
process in the ready
queue



OS (scheduler)

Lock



Disk

