

# Operating Systems Principles

## CPU Management (Processes)

# Choosing the Right Scheduling Algorithm/Scheduling Criteria

- Performance
  - Waiting time or response time
    - Response time – amount of time it takes from when a request was submitted until the first response is produced
    - Worth reducing in time-sharing environments
  - Throughput – # of processes that complete their execution per time unit
    - Worth increasing in batch-processing environments
- Fairness
  - E.g., proportional sharing
  - E.g., reservations

# What's Wrong with FCFS?

- First-Come-First-Served
  - A natural and easy to understand/implement policy
  - What's wrong with it?
- Lets consider some better options
  - Also pay attention to data structures that the scheduler might want to use for each
  - Efficiency of the scheduler itself – time/space complexity

## First- Come, First-Served (FCFS) Scheduling

<u>Process</u>	<u>Burst Time</u>
$P_1$	24
$P_2$	3
$P_3$	3

- Suppose that the processes arrive in the order:  $P_1, P_2, P_3$   
The Gantt Chart for the schedule is:



- Waiting time for  $P_1 = 0$ ;  $P_2 = 24$ ;  $P_3 = 27$
- Average waiting time:  $(0 + 24 + 27)/3 = 17$

# FCFS Scheduling (Cont.)

Suppose that the processes arrive in the order:

$$P_2, P_3, P_1$$

- The Gantt chart for the schedule is:



- Waiting time for  $P_1 = 6$ ;  $P_2 = 0$ ;  $P_3 = 3$
- Average waiting time:  $(6 + 0 + 3)/3 = 3$
- Much better than previous case
- **Convoy effect** - short process behind long process
  - Consider one CPU-bound and many I/O-bound processes


# Shortest-Job-First (SJF)

- Associate with each process the length of its next CPU burst. Use these lengths to schedule the process with the shortest time
- SJF is optimal for avg. waiting time – gives minimum average waiting time for a given set of processes
  - In class: Compute average waiting time for an example with SJF and compare w/ FCFS
  - Exercise: Prove the optimality claimed above

# Example of SJF

<u>Process</u>	<u>Burst Time</u>
$P_1$	6
$P_2$	8
$P_3$	7
$P_4$	3

- SJF scheduling chart



- Average waiting time =  $(3 + 16 + 9 + 0) / 4 = 7$

# Why Pre-emption is Necessary

- To improve CPU utilization
  - Most processes are not *ready* at all times during their lifetimes
  - E.g., think of a text editor waiting for input from the keyboard
  - Also improves I/O utilization
- To improve responsiveness
  - Many processes would prefer “slow but steady progress” over “long wait followed by fast process”
- Most modern CPU schedulers are pre-emptive



# SJF: Variations on the theme


- **Non-preemptive:** once CPU given to the process it cannot be preempted until completes its CPU burst - the SJF we already saw
- **Preemptive:** if a new process arrives with CPU length less than remaining time of current executing process, preempt. This scheme is known as *Shortest-Remaining-Time-First (SRTF)*
  - Also called *Shortest Remaining Processing Time (SRPT)*
- Why SJF/SRTF may not be practical
  - CPU requirement of a process rarely known in advance

# Example of Shortest-remaining-time-first

- Now we add the concepts of varying arrival times and preemption to the analysis

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0	8
$P_2$	1	4
$P_3$	2	9
$P_4$	3	5

- Preemptive* SJF Gantt Chart



- Average waiting time =  $[(10-1)+(1-1)+(17-2)+5-3]/4 = 26/4 = 6.5$  msec