

# Operating Systems Principles

## CPU Management

## CPU Scheduling (2), Process Synchronization


# Round Robin (RR)

- Each process gets a small unit of CPU time (*time quantum*), usually 10-100 milliseconds. After this time has elapsed, the process is preempted and added to the end of the ready queue.
- If there are  $n$  processes in the ready queue and the time quantum is  $q$ , then each process gets  $1/n$  of the CPU time in chunks of at most  $q$  time units at once. No process waits more than  $(n-1)q$  time units.
- Performance
  - $q$  large  $\Rightarrow$  FCFS
  - $q$  small  $\Rightarrow q$  must be large with respect to context switch, otherwise overhead is too high

## Example of RR with Time Quantum = 20


<u>Process</u>	<u>CPU Time</u>
$P_1$	53
$P_2$	17
$P_3$	68
$P_4$	24

- The Gantt chart is:




- Typically, higher average turnaround than SJF, but better *response*

# Time Quantum and Context Switch Time



# Turnaround Time Varies With Time Quantum




# Proportional-Share Schedulers

- A generalization of round robin
- Process  $P_i$  given a CPU weight  $w_i > 0$
- The scheduler needs to ensure the following
  - *For all*  $i, j$ ,  $|T_i(t_1, t_2)/T_j(t_1, t_2) - w_i/w_j| \leq \epsilon$
  - Given  $P_i$  and  $P_j$  were backlogged during  $[t_1, t_2]$
- Que: Who chooses the weights and how?

# Lottery Scheduling

- Perhaps the simplest proportional-share scheduler
- Create lottery tickets equal to the sum of the weights of all processes
- Draw a lottery ticket and schedule the process that owns that ticket

# Lottery Scheduling Example



9

$P1=6$

1	4
2	5
3	6


$P2=9$

7	10	13
8	11	14
9	12	15

*Schedule P2*



# Lottery Scheduling Example



3


$P1=6$

$P2=9$

1	4	7	10	13
2	5	8	11	14
3	6	9	12	15

*Schedule P1*

# Lottery Scheduling Example



11

$P1=6$

1	4
2	5
3	6

$P2=9$

7	10	13
8	11	14
9	12	15

- As  $t \rightarrow \infty$ , processes will get their share (unless they were blocked a lot)
- Problem with Lottery scheduling: Only probabilistic guarantee
- What does the scheduler have to do
  - When a new process arrives?
  - When a process terminates?

*Schedule P2*

# Lottery Scheduling

- Exercise: Calculate the time complexity of the operations Lottery scheduling will involve

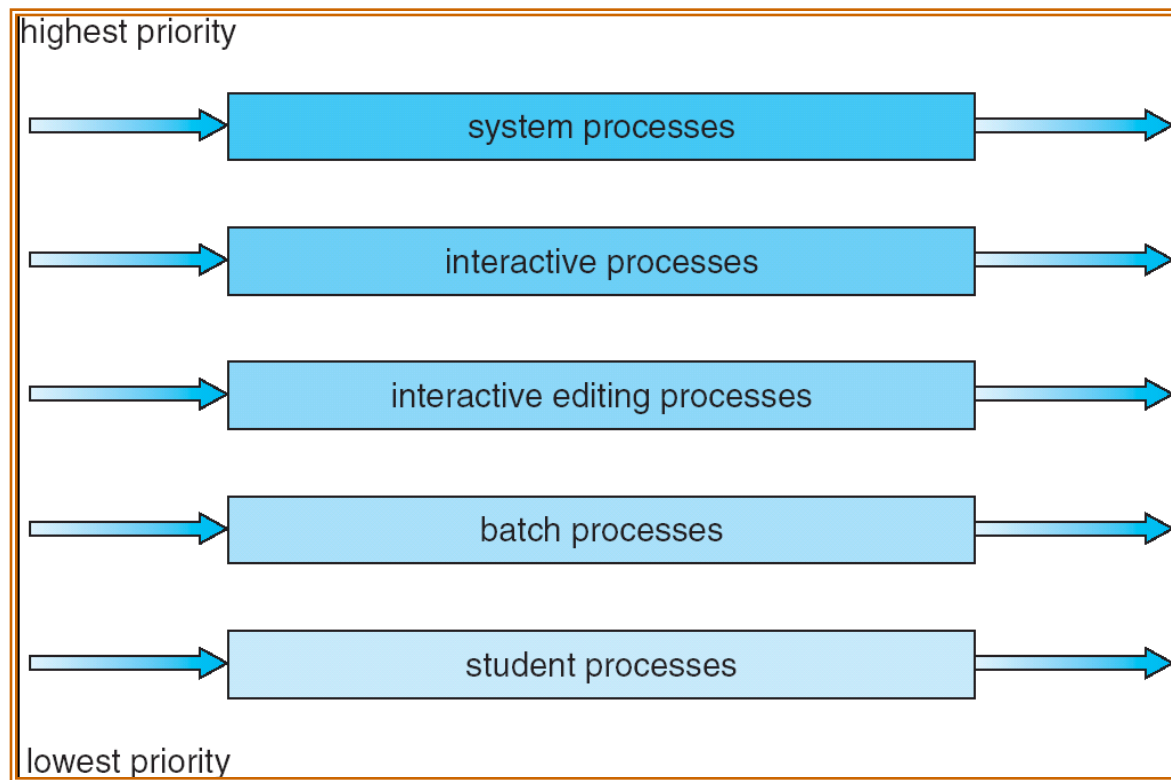
# Priority-based Scheduling

- Associate with each process a quantity called its *CPU priority*
- At each scheduling instant
  - Pick the ready process with the highest CPU priority
  - Update (usually decrement) the priority of the process last running
    - Priority = Time since arrival => FCFS
    - Priority =  $1/\text{Size}$  => SJF
    - Priority =  $1/\text{Remaining Time}$  => SRPT
    - Priority = Time since last run => Round-robin (RR)
- UNIX variants
  - Priority values are positive integers with upper bounds
  - Decreased every quantum
    - Fairness, avoid starvation
  - Increased if the process was waiting, more wait => larger increase
    - To make interactive processes more responsive
  - Problems
    - Hard to analyze theoretically, so hard to give any guarantees
    - May unfairly reward blocking processes

# Multilevel Queue

- Ready queue is partitioned into separate queues:
  - foreground (interactive)
  - background (batch)
- Each queue has its own scheduling algorithm
  - foreground – RR
  - background – FCFS
- Scheduling must be done between the queues
  - Fixed priority scheduling; (i.e., serve all from foreground then from background).
    - Possibility of starvation.
  - Time slice – each queue gets a certain amount of CPU time which it can schedule amongst its processes;
    - i.e., 80% to foreground in RR, 20% to background in FCFS

# Multilevel Queue Scheduling



# Multilevel Feedback Queue


- A process can move between the various queues; aging can be implemented this way
- Multilevel-feedback-queue scheduler defined by the following parameters:
  - number of queues
  - scheduling algorithms for each queue
  - method used to determine when to upgrade a process
  - method used to determine when to demote a process
  - method used to determine which queue a process will enter when that process needs service

# Example of Multilevel Feedback Queue

- Three queues:
  - $Q_0$  – RR with time quantum 8 milliseconds
  - $Q_1$  – RR time quantum 16 milliseconds
  - $Q_2$  – FCFS
- Scheduling
  - A new job enters queue  $Q_0$  which is served FCFS. When it gains CPU, job receives 8 milliseconds. If it does not finish in 8 milliseconds, job is moved to queue  $Q_1$ .
  - At  $Q_1$  job is again served FCFS and receives 16 additional milliseconds. If it still does not complete, it is preempted and moved to queue  $Q_2$ .



# Multilevel Feedback Queues




# Work Conservation

- Examples of work-conserving schedulers: All schedulers we have studied so far
  - Work conservation: The resource (e.g., CPU) can not be idle if there is some work to be done (e.g., at least one ready process)
- Example of a non-work-conserving scheduler:
  - Reservation-based (coming up)

# Reservation-based Schedulers


- Each process has a pair  $(x, y)$ 
  - Divide time into periods of length  $y$  each
  - Guaranteed to get  $x$  time units every period
- Why is this NWC?
- Why design a NWC scheduler?

# Hierarchical Schedulers



- A way to compose a variety of schedulers
- Sets of processes with different scheduling needs
- Exercise: Think of conditions under which a tree of schedulers becomes NWC

# Address Spaces: A Refinement




- The OS is made part of the address space of every process
- Why?
  - User/kernel transitions are not context switches anymore!
  - Need to disallow virtual address range of OS from being accessed by processes -> MMU helps with this

# Process Synchronization

- We will study it in the context of the shared memory model

# Motivating Example: Producer and Consumer



```
while (true) {  
    /* produce an item and put in nextProduced */  
    while (count == BUFFER_SIZE); // do nothing  
  
    buffer [in] = nextProduced;  
    in = (in + 1) % BUFFER_SIZE;  
    count++;  
}
```

```
while (true) {  
    while (count == 0); // do nothing  
  
    nextConsumed = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    count--;  
    /* consume the item in nextConsumed  
}
```

- Shared variables: buffer and count
- Local variables: in, out, nextProduced, nextConsumed